

Deep Reinforcement Learning by S. Levine

CS285 - UC Berkeley

Filippo GIRUZZI

January 21, 2020

Contents

1	Introduction	4
2	Supervised Learning of behaviors	4
2.1	Goal	4
2.2	Algorithms	4
2.2.1	DAGger: Dataset Aggregation	4
2.3	Tips & hacks	4
3	Introduction to Reinforcement Learning	5
3.1	Goal	5
3.2	Algorithms	5
3.2.1	Global structure	5
3.2.2	Exemples	6
4	Policy Gradients	6
4.1	Algorithms	6
4.1.1	REINFORCE	6
4.2	Tips & hacks	6
5	Actor-Critic algorithms	7
5.1	Algorithms	7
5.1.1	Batch Actor-Critic	7
5.1.2	Online Actor-Critic	7
5.2	Tips & hacks	7
6	Value function methods	8
6.1	Algorithms	8
6.1.1	Policy iteration	8
6.1.2	Policy iteration with Dynamic programming	8
6.1.3	Value iteration	8
6.1.4	Fitted Value iteration	9
6.1.5	Fitted Q-iteration	9
6.1.6	Online Q-iteration	9
6.2	Tips & hacks	9
7	Deep Reinforcement Learning with Q-functions	10
7.1	Algorithms	10
7.1.1	Q-learning with replay buffer	10
7.1.2	Q-learning with replay buffer and target network	10
7.1.3	DQN: classic Deep Q-learning	10
7.1.4	DDPG: Q-learning for continuous actions	11
7.2	Tips & hacks	11
8	Advanced Policy Gradients	12
9	Model-based planning	12

10 Model-based Reinforcement Learning	12
10.1 Algorithms	12
10.1.1 Model-based Reinforcement Learning version 0.5	12
10.1.2 Model-based Reinforcement Learning version 1.0	12
10.1.3 Model-based Reinforcement Learning version 1.5	12
10.1.4 Model-based Reinforcement Learning with latent space models . . .	13
10.2 Tips & hacks	13
11 Model-based Policy Learning	13
11.1 Algorithms	13
11.1.1 Model-based Reinforcement Learning version 2.0	13
11.1.2 DYNA: online Q-learning model-free Reinforcement Learning with a model	14
11.1.3 General DYNA-style model-based Reinforcement Learning	14
11.1.4 MBA: Model-based Acceleration – MVE: Model-based Value Ex- pansion – MBPO: Model-based Policy Optimization	14
11.1.5 Divide and Conquer Reinforcement Learning	15
11.2 Tips & hacks	15
12 Variational Inference & Generative models	15
13 Control as inference	15
13.1 Algorithms	15
13.1.1 Soft Q-learning	15
14 Inverse Reinforcement Learning	16
14.1 Algorithms	16
14.1.1 Maximum Entropy Inverse Reinforcement Learning	16
14.2 Tips & hacks	16
15 Transfer & Multi-task Learning	16
15.1 Tips & hacks	16
16 Distributed Reinforcement Learning	17
17 Exploration	17
17.1 Algorithms	17
17.1.1 Pre-train & finetune	17
17.2 Tips & hacks	17
18 Meta-learning	18
19 Information theory	18

1 Introduction

2 Supervised Learning of behaviors

2.1 Goal

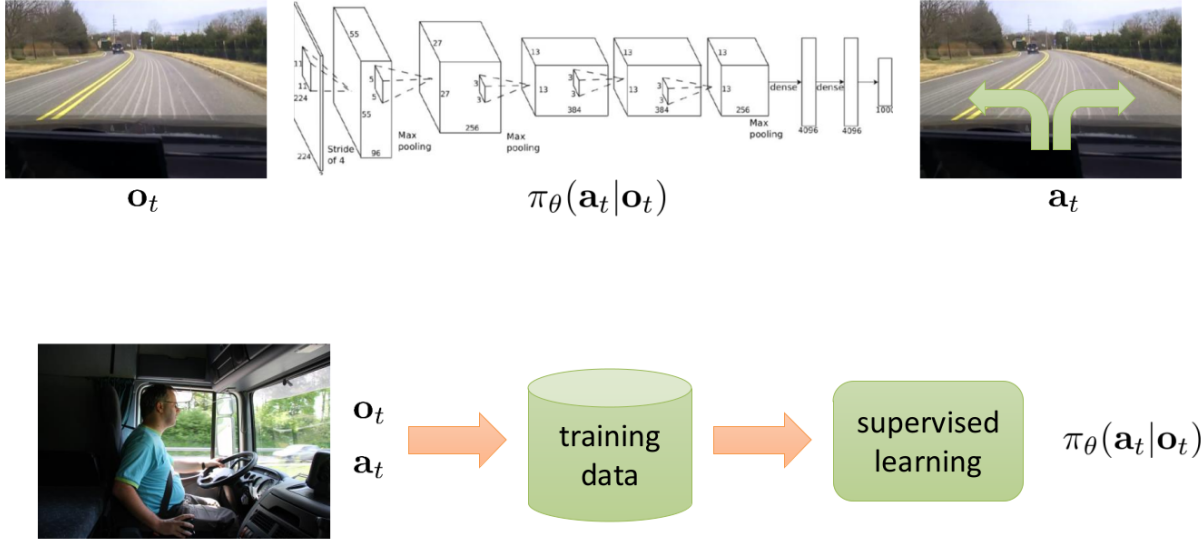


Figure 1: Imitation Learning objective

2.2 Algorithms

2.2.1 DAgger: Dataset Aggregation

Algorithm 1: DAgger

Loop

1. Train $\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_N, \mathbf{a}_N\}$
2. Run $\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
3. Ask human to label \mathcal{D}_π with actions \mathbf{a}_t
4. Aggregate $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

end

2.3 Tips & hacks

1. Distributional shift:

- Use DAgger
- Randomize input

2. Non-Markovian behavior: use RNN to model $\pi_\theta(\mathbf{a}_t | \mathbf{o}_1, \dots, \mathbf{o}_t)$

3. Multi-modal behavior:

- Output a mixture of Gaussians
- Use latent variable models
- Use autoregressive discretization

3 Introduction to Reinforcement Learning

3.1 Goal

Definition 1 (POMDP: Partially Observed Markov Process)

$$\mathcal{M} = \{\mathcal{S}, \mathcal{A}, \mathcal{O}, \mathcal{T}, \mathcal{E}, r\} \quad (1)$$

Definition 2 (Reinforcement Learning objective)

$$\max_{\theta} E_{(\mathbf{s}, \mathbf{a}) \sim p_{\theta}(\mathbf{s}, \mathbf{a})} [r(\mathbf{s}, \mathbf{a})] \quad (2)$$

$$\max_{\theta} \sum_{t=1}^T E_{(\mathbf{s}_t, \mathbf{a}_t) \sim p_{\theta}(\mathbf{s}_t, \mathbf{a}_t)} [r(\mathbf{s}_t, \mathbf{a}_t)] = \max_{\theta} E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right] \quad (3)$$

Definition 3 (Q-function)

$$Q^{\pi}(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^T E_{\pi_{\theta}} [r(\mathbf{s}'_{t'}, \mathbf{a}'_{t'}) | \mathbf{s}_t, \mathbf{a}_t] \quad (4)$$

Definition 4 (Value function)

$$V^{\pi}(\mathbf{s}_t) = \sum_{t'=t}^T E_{\pi_{\theta}} [r(\mathbf{s}'_{t'}, \mathbf{a}'_{t'}) | \mathbf{s}_t] = E_{\mathbf{a}_t \sim \pi(\mathbf{a}_t | \mathbf{s}_t)} [Q^{\pi}(\mathbf{s}_t, \mathbf{a}_t)] \quad (5)$$

$E_{\mathbf{s}_1 \sim p(\mathbf{s}_1)} [V^{\pi}(\mathbf{s}_1)]$ is the Reinforcement Learning objective.

3.2 Algorithms

3.2.1 Global structure

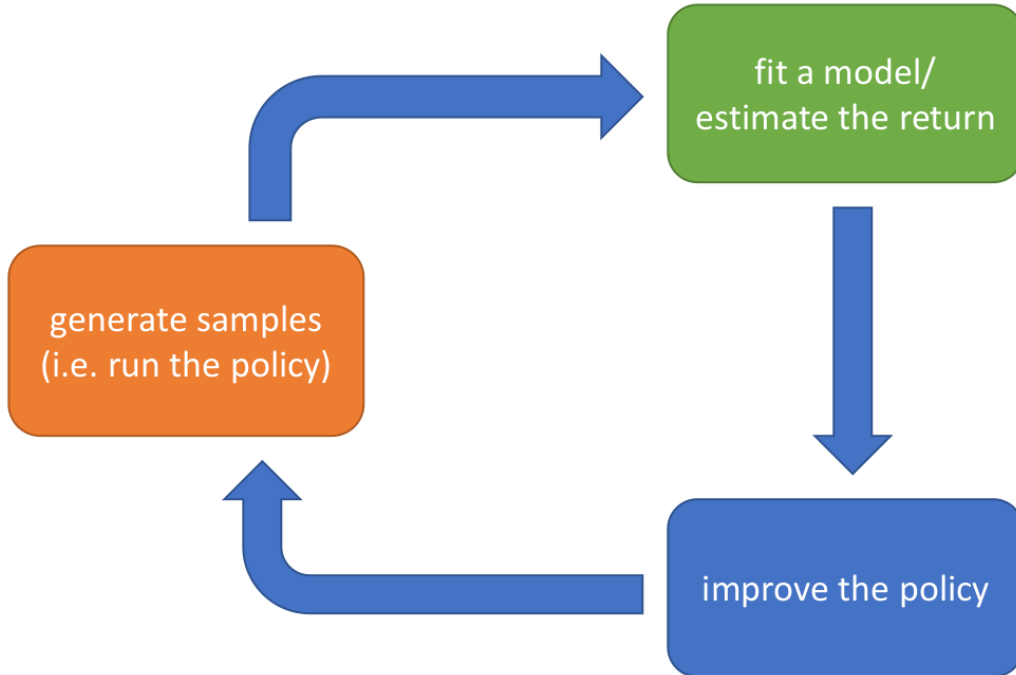


Figure 2: Global RL algorithm structure

3.2.2 Examples

1. **Policy Gradients:** directly differentiate the RL objective
 - REINFORCE
 - Natural Policy Gradient
 - Trust region policy optimization
2. **Value-based:** estimate the Value function or the Q-function of the optimal policy
 - Q-learning, DQN
 - Temporal difference learning
 - Fitted Value iteration
3. **Actor-critic:** estimate the Value function or the Q-function of the current policy, and use it to improve the policy
 - A3C: Asynchronous advantage Actor-critic
 - SAC: Soft Actor-critic
4. **Model-based:** estimate a transition model, and use it for planning, to improve a policy or else
 - DYNA
 - Guided policy search

4 Policy Gradients

4.1 Algorithms

4.1.1 REINFORCE

Algorithm 2: REINFORCE

```
Loop
  1. Run the policy and sample  $\{\tau^i\}$  from  $\pi_\theta(\mathbf{a}_t|\mathbf{s}_t)$ 
  2.  $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left( \sum_{t=1}^T \nabla_\theta \log \pi_\theta(\mathbf{a}_t^i|\mathbf{s}_t^i) \right) \left( \sum_{t=1}^T r(\mathbf{s}_t^i, \mathbf{a}_t^i) \right)$ 
  3.  $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$ 
end
```

4.2 Tips & hacks

1. Variance reduction:

- Causality & reward to-go: $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_\theta \log \pi_\theta(\mathbf{a}_t^i|\mathbf{s}_t^i) \right) \hat{Q}_t^i$ where

$$\text{the reward to-go is } \hat{Q}_t^i = \sum_{t'=t}^T r(\mathbf{s}_{t'}^i, \mathbf{a}_{t'}^i)$$

- Use baselines: $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(\tau)(r(\tau) - b)$ where the baseline can be $b = \frac{1}{N} \sum_{i=1}^N r(\tau)$

2. **Off-policy Policy Gradient:** use importance sampling

3. **Discount factor:** $\hat{Q}_t^i = \sum_{t'=t}^T \gamma^{t'-t} r(\mathbf{s}_{t'}^i, \mathbf{a}_{t'}^i)$

5 Actor-Critic algorithms

5.1 Algorithms

5.1.1 Batch Actor-Critic

Algorithm 3: Batch Actor-Critic

Loop

1. Run the policy and sample $\{\mathbf{s}_i, \mathbf{a}_i\}$ from $\pi_{\theta}(\mathbf{a}|\mathbf{s})$
2. Fit $\hat{V}_{\phi}^{\pi}(\mathbf{s})$ to $y_t^i = \sum_{t'=t}^T \gamma^{t'-t} r(\mathbf{s}_{t'}^i, \mathbf{a}_{t'}^i) = r(\mathbf{s}_t^i, \mathbf{a}_t^i) + \gamma \hat{V}_{\phi}^{\pi}(\mathbf{s}_{t+1}^i)$
3. Evaluate $\hat{A}^{\pi}(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_{\phi}^{\pi}(\mathbf{s}_i') - \hat{V}_{\phi}^{\pi}(\mathbf{s}_i)$
4. $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_i|\mathbf{s}_i) \hat{A}^{\pi}(\mathbf{s}_i, \mathbf{a}_i)$
5. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

end

5.1.2 Online Actor-Critic

Algorithm 4: Online Actor-Critic

Loop

1. Take an action $\mathbf{a} \sim \pi_{\theta}(\mathbf{a}|\mathbf{s})$ and get $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$
2. Update \hat{V}_{ϕ}^{π} using target $r + \gamma \hat{V}_{\phi}^{\pi}(\mathbf{s}')$
3. Evaluate $\hat{A}^{\pi}(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \hat{V}_{\phi}^{\pi}(\mathbf{s}') - \hat{V}_{\phi}^{\pi}(\mathbf{s})$
4. $\nabla_{\theta} J(\theta) \approx \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}|\mathbf{s}) \hat{A}^{\pi}(\mathbf{s}, \mathbf{a})$
5. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

end

5.2 Tips & hacks

1. **Architecture design:** use two networks or a single MTL network to fit $\pi_{\theta}(\mathbf{a}|\mathbf{s})$ and $\hat{V}_{\phi}^{\pi}(\mathbf{s})$
2. **Parallel workers:** use parallel workers to batch steps **2.** and **4.** of the Online algorithm, either synchronized or asynchronous

3. **n-step returns:** $\hat{A}_n^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) + \gamma^n \hat{V}_\phi^\pi(\mathbf{s}_{t+n}) - \hat{V}_\phi^\pi(\mathbf{s}_t)$ with $n > 1$
4. **Generalized Advantage Estimation:** $\hat{A}_{GAE}^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{n=1}^{\infty} w_n \hat{A}_n^\pi(\mathbf{s}_t, \mathbf{a}_t)$ with $w_n \propto \lambda^{n-1}$ such that $\hat{A}_{GAE}^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^{\infty} (\gamma\lambda)^{t'-t} \delta_{t'}$ where $\delta_{t'} = r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) + \gamma \hat{V}_\phi^\pi(\mathbf{s}_{t'+1}) - \hat{V}_\phi^\pi(\mathbf{s}_{t'})$

6 Value function methods

6.1 Algorithms

6.1.1 Policy iteration

Algorithm 5: Policy iteration

```

Loop
  1. Evaluate  $A^\pi(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma E[V^\pi(\mathbf{s}')] - V^\pi(\mathbf{s})$ 
  2.  $\pi \leftarrow \pi'$  where  $\pi'(\mathbf{a}_t | \mathbf{s}_t) = \begin{cases} 1 & \text{if } \mathbf{a}_t = \arg\max_{\mathbf{a}_t} A^\pi(\mathbf{s}_t, \mathbf{a}_t) \\ 0 & \text{otherwise} \end{cases}$ 
end

```

6.1.2 Policy iteration with Dynamic programming

Algorithm 6: Policy iteration with Dynamic programming

```

Loop
  1. Evaluate  $V^\pi(\mathbf{s}) = r(\mathbf{s}, \pi(\mathbf{s})) + \gamma E_{\mathbf{s}' \sim p(\mathbf{s}' | \mathbf{s}, \pi(\mathbf{s}))} [V^\pi(\mathbf{s}')]$ 
  2.  $\pi \leftarrow \pi'$  where  $\pi'(\mathbf{a}_t | \mathbf{s}_t) = \begin{cases} 1 & \text{if } \mathbf{a}_t = \arg\max_{\mathbf{a}_t} A^\pi(\mathbf{s}_t, \mathbf{a}_t) \\ 0 & \text{otherwise} \end{cases}$ 
end

```

6.1.3 Value iteration

Algorithm 7: Value iteration

```

Loop
  1. Set  $Q^\pi(\mathbf{s}, \mathbf{a}) \leftarrow r(\mathbf{s}, \mathbf{a}) + \gamma E_{\mathbf{s}' \sim p(\mathbf{s}' | \mathbf{s}, \mathbf{a})} [V^\pi(\mathbf{s}')]$ 
  2. Set  $V^\pi(\mathbf{s}) \leftarrow \max_{\mathbf{a}} Q^\pi(\mathbf{s}, \mathbf{a})$ 
end

```

6.1.4 Fitted Value iteration

Algorithm 8: Fitted Value iteration

```

Loop
  1. Collect dataset  $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$  using some policy
  Loop
    2. Set  $y_i \leftarrow \max_{\mathbf{a}_i} (r(\mathbf{s}_i, \mathbf{a}_i) + \gamma E[V_\phi(\mathbf{s}'_i)])$ 
    3. Set  $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|V_\phi(\mathbf{s}_i) - y_i\|^2$ 
  end
end

```

6.1.5 Fitted Q-iteration

Algorithm 9: Fitted Q-iteration

```

Loop
  1. Collect dataset  $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$  using some policy
  Loop
    2. Set  $y_i \leftarrow r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'_i} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$ 
    3. Set  $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - y_i\|^2$ 
  end
end

```

6.1.6 Online Q-iteration

Algorithm 10: Online Q-iteration

```

Loop
  1. Take some action  $\mathbf{a}_i$  and observe  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$ 
  2.  $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'_i} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$ 
  3.  $\phi \leftarrow \phi - \alpha \frac{dQ_\phi}{d\phi}(\mathbf{s}_i, \mathbf{a}_i)(Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - y_i)$ 
end

```

6.2 Tips & hacks

1. **Exploration:** Use a different policy to improve exploration

- Epsilon-greedy exploration: $\pi(\mathbf{a}_t | \mathbf{s}_t) = \begin{cases} 1 - \epsilon & \text{if } \mathbf{a}_t = \operatorname{argmax}_{\mathbf{a}_t} Q_\phi(\mathbf{s}_t, \mathbf{a}_t) \\ \frac{\epsilon}{|\mathcal{A}| - 1} & \text{otherwise} \end{cases}$
- Boltzmann exploration: $\pi(\mathbf{a}_t | \mathbf{s}_t) \propto \exp(Q_\phi(\mathbf{s}_t, \mathbf{a}_t))$

7 Deep Reinforcement Learning with Q-functions

7.1 Algorithms

7.1.1 Q-learning with replay buffer

Algorithm 11: Q-learning with replay buffer

```
Loop
  1. Collect dataset  $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$  using some policy, and add it to  $\mathcal{B}$ 
  Loop
    2. Sample a batch  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$  from  $\mathcal{B}$ 
    3.  $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'_i} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$ 
    4.  $\phi \leftarrow \phi - \alpha \sum_i \frac{dQ_\phi}{d\phi}(\mathbf{s}_i, \mathbf{a}_i)(Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - y_i)$ 
  end
end
```

7.1.2 Q-learning with replay buffer and target network

Algorithm 12: Q-learning with replay buffer

```
Loop
  1. Save network parameters  $\phi' \leftarrow \phi$ 
  Loop
    2. Collect dataset  $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$  using some policy, and add it to  $\mathcal{B}$ 
    Loop
      3. Sample a batch  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$  from  $\mathcal{B}$ 
      4.  $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'_i} Q_{\phi'}(\mathbf{s}'_i, \mathbf{a}'_i)$ 
      5.  $\phi \leftarrow \phi - \alpha \sum_i \frac{dQ_\phi}{d\phi}(\mathbf{s}_i, \mathbf{a}_i)(Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - y_i)$ 
    end
  end
end
```

7.1.3 DQN: classic Deep Q-learning

Algorithm 13: DQN: classic Deep Q-learning

```
Loop
  1. Take some action  $\mathbf{a}_i$ , observe  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$  and add it to  $\mathcal{B}$ 
  2. Sample a mini-batch  $(\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j)$  from  $\mathcal{B}$  uniformly
  3. Compute  $y_j = r(\mathbf{s}_j, \mathbf{a}_j) + \gamma \max_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mathbf{a}'_j)$  using the target network  $Q_{\phi'}$ 
  4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_\phi(\mathbf{s}_j, \mathbf{a}_j) - y_j)$ 
  5. Update  $\phi'$ : copy  $\phi$  every  $N$  steps or using Polyak averaging
      $\phi' \leftarrow \tau \phi' + (1 - \tau)\phi$ 
end
```

7.1.4 DDPG: Q-learning for continuous actions

Algorithm 14: DDPG: Deep Deterministic Policy Gradient

Loop

1. Take some action \mathbf{a}_i , observe $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$ and add it to \mathcal{B}
2. Sample a mini-batch $(\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j)$ from \mathcal{B} uniformly
3. Compute $y_j = r(\mathbf{s}_j, \mathbf{a}_j) + \gamma \max_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mu_{\theta'}(\mathbf{s}'_j))$ using target nets $Q_{\phi'}, \mu_{\theta'}$
4. $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_\phi(\mathbf{s}_j, \mathbf{a}_j) - y_j)$
5. $\theta \leftarrow \theta + \beta \sum_j \frac{d\mu}{d\theta}(\mathbf{s}_j) \frac{dQ_\phi}{d\mathbf{a}}(\mathbf{s}_j, \mathbf{a})$
6. Update ϕ' and θ'

end

7.2 Tips & hacks

1. **Double Q-learning:** use two networks to avoid overestimation

- Standard Q-learning: $y = r + \gamma Q_{\phi'}(\mathbf{s}', \arg\max_{\mathbf{a}'} Q_{\phi'}(\mathbf{s}', \mathbf{a}'))$
- Double Q-learning: $y = r + \gamma Q_{\phi'}(\mathbf{s}', \arg\max_{\mathbf{a}'} Q_{\phi}(\mathbf{s}', \mathbf{a}'))$

2. **Multi-step returns:** $y_t^j = \sum_{t'=t}^{t+N-1} \gamma^{t'-t} r_{t'}^j + \gamma^N \max_{\mathbf{a}_{t+N}^j} Q_{\phi'}(\mathbf{s}_{t+N}^j, \mathbf{a}_{t+N}^j)$

3. **General tips:**

- Test on easy, reliable tasks
- Use large replay buffers
- Start with high exploration, then gradually reduce
- Clip gradients or use Huber loss
- Use double Q-learning
- Use N-step returns
- Schedule exploration and learning rates
- Run multiple random seeds
- Use DDPG for continuous actions, or other methods

8 Advanced Policy Gradients

9 Model-based planning

10 Model-based Reinforcement Learning

10.1 Algorithms

10.1.1 Model-based Reinforcement Learning version 0.5

Algorithm 15: Model-based RL version 0.5

1. Run some base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$, and collect $\mathcal{D} = \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
 2. Learn the dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
 3. Plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions
-

10.1.2 Model-based Reinforcement Learning version 1.0

Algorithm 16: Model-based RL version 1.0

1. Run some base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$, and collect $\mathcal{D} = \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
- Loop**
2. Learn the dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
 3. Plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions
 4. Execute those actions and add the data to \mathcal{D}
- end**
-

10.1.3 Model-based Reinforcement Learning version 1.5

Algorithm 17: Model-based RL version 1.5

1. Run some base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$, and collect $\mathcal{D} = \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
- Loop** (*every N steps*)
2. Learn the dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
- Loop**
3. Plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions
 4. Execute the first planned action, and observe \mathbf{s}' (MPC)
 5. Append $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ to \mathcal{D}
- end**
- end**
-

10.1.4 Model-based Reinforcement Learning with latent space models

Algorithm 18: Model-based RL with latent space models

```

1. Run some base policy  $\pi_0(\mathbf{a}_t|\mathbf{o}_t)$ , and collect  $\mathcal{D} = \{(\mathbf{o}_i, \mathbf{a}_i, \mathbf{o}'_i)\}$ 
Loop (every  $N$  steps)
    2. Learn  $p_\phi(\mathbf{s}_{t+1}|\mathbf{s}_t, \mathbf{a}_t)$ ,  $p_\phi(\mathbf{r}_t|\mathbf{s}_t)$ ,  $p(\mathbf{o}_t|\mathbf{s}_t)$ ,  $g_\psi(\mathbf{o}_t)$  to maximize
        
$$\max_{\phi, \psi} \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \underbrace{\log p_\phi(g_\psi(\mathbf{o}'_{t+1})|g_\psi(\mathbf{o}_t), \mathbf{a}_t^i)}_{\text{latent space dynamics}} + \underbrace{\log p_\phi(\mathbf{o}_t^i|g_\psi(\mathbf{o}_t^i))}_{\text{image reconstruction}} + \underbrace{\log p_\phi(r_t^i|g_\psi(\mathbf{o}_t^i))}_{\text{reward model}}$$

        Loop
            3. Plan through the model to choose actions
            4. Execute the first planned action, and observe  $\mathbf{o}'$ 
            5. Append  $(\mathbf{o}, \mathbf{a}, \mathbf{o}')$  to  $\mathcal{D}$ 
        end
    end
end

```

10.2 Tips & hacks

1. Uncertainty-aware models:

- Use output entropy in the networks
- Estimate the uncertainty of the networks
 - Use Bayesian networks
 - Use bootstrap ensembles, where each model is trained on a subset sampled with replacement from the dataset

11 Model-based Policy Learning

11.1 Algorithms

11.1.1 Model-based Reinforcement Learning version 2.0

Algorithm 19: Model-based RL version 2.0

```

1. Run some base policy  $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$ , and collect  $\mathcal{D} = \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$ 
Loop
    2. Learn the dynamics model  $f(\mathbf{s}, \mathbf{a})$  to minimize  $\sum_i ||f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i||^2$ 
    3. Backpropagate through  $f(\mathbf{s}, \mathbf{a})$  into the policy to optimize  $\pi_\theta(\mathbf{a}_t|\mathbf{s}_t)$ 
    4. Run  $\pi_\theta(\mathbf{a}_t|\mathbf{s}_t)$  and add the visited  $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$  to  $\mathcal{D}$ 
end

```

11.1.2 DYNA: online Q-learning model-free Reinforcement Learning with a model

Algorithm 20: DYNA

```

Loop
  1. Given  $\mathbf{s}$ , pick  $\mathbf{a}$  using some exploration policy, and observe  $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$ 
  2. Update the models  $\hat{p}(\mathbf{s}'|\mathbf{s}, \mathbf{a})$  and  $\hat{r}(\mathbf{s}, \mathbf{a})$  using  $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ 
  3. Update  $Q(\mathbf{s}, \mathbf{a}) \leftarrow Q(\mathbf{s}, \mathbf{a}) + \alpha E_{\mathbf{s}', r} [r + \max_{\mathbf{a}'} Q(\mathbf{s}', \mathbf{a}') - Q(\mathbf{s}, \mathbf{a})]$ 
  Loop
    4. Sample  $(\mathbf{s}, \mathbf{a}) \sim \mathcal{B}$  from the buffer
    5. Update  $Q(\mathbf{s}, \mathbf{a}) \leftarrow Q(\mathbf{s}, \mathbf{a}) + \alpha E_{\mathbf{s}', r} [r + \max_{\mathbf{a}'} Q(\mathbf{s}', \mathbf{a}') - Q(\mathbf{s}, \mathbf{a})]$ 
  end
end

```

11.1.3 General DYNA-style model-based Reinforcement Learning

Algorithm 21: General DYNA-style model-based RL

```

Loop
  1. Collect some data  $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$ 
  2. Learn the model  $\hat{p}(\mathbf{s}'|\mathbf{s}, \mathbf{a})$ , and optionally  $\hat{r}(\mathbf{s}, \mathbf{a})$ 
  Loop
    3. Sample  $\mathbf{s} \sim \mathcal{B}$  from the buffer
    4. Choose action  $\mathbf{a}$  (from  $\mathcal{B}$ , from  $\pi$  or randomly)
    5. Simulate  $\mathbf{s}' \sim \hat{p}(\mathbf{s}'|\mathbf{s}, \mathbf{a})$ , and  $r = \hat{r}(\mathbf{s}, \mathbf{a})$ 
    6. Train on  $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$  with model-free RL
    7. (Optional) Take  $N$  more model-based steps
  end
end

```

11.1.4 MBA: Model-based Acceleration – MVE: Model-based Value Expansion – MBPO: Model-based Policy Optimization

Algorithm 22: MBA – MVE – MBPO

```

Loop
  1. Take some action  $\mathbf{a}_i$ , observe  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$  and add it to  $\mathcal{B}$ 
  2. Sample a mini-batch  $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j\}$  from  $\mathcal{B}$  uniformly
  3. Use  $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j\}$  to update the model  $\hat{p}(\mathbf{s}'|\mathbf{s}, \mathbf{a})$ 
  4. Sample  $\{\mathbf{s}_j\}$  from  $\mathcal{B}$ 
  5. For each  $\mathbf{s}_j$ , perform model-based rollout with  $\mathbf{a} = \pi(\mathbf{s})$ 
  6. Use all transitions  $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$  along the rollout to update the Q-function
end

```

11.1.5 Divide and Conquer Reinforcement Learning

Algorithm 23: Divide and Conquer RL

Loop

1. Optimize each local policy $\pi_{\theta_i}(\mathbf{a}_t|\mathbf{s}_t)$ on initial state \mathbf{s}_0^i w.r.t $\tilde{r}_k^i(\mathbf{s}_t, \mathbf{a}_t)$
2. Use samples from step 1. to train $\pi_{\theta}(\mathbf{a}_t|\mathbf{s}_t)$ to mimic each $\pi_{\theta_i}(\mathbf{a}_t|\mathbf{s}_t)$
3. Update reward $\tilde{r}_{k+1}^i(\mathbf{s}_t, \mathbf{a}_t) = r(\mathbf{s}_t, \mathbf{a}_t) + \lambda_{k+1}^i \log \pi_{\theta}(\mathbf{a}_t|\mathbf{s}_t)$

end

11.2 Tips & hacks

1. **Local models:** use complex controllers as local models
2. **Dynamics models:** use Bayesian linear regression
3. **Multi-task learning:**
 - Use ensemble models
 - Use soft targets
 - Use distillation: train a global policy with supervised learning to mimic each policy

12 Variational Inference & Generative models

13 Control as inference

13.1 Algorithms

13.1.1 Soft Q-learning

Algorithm 24: Soft Q-learning

Loop

1. Take some action \mathbf{a}_i , observe $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$ and add it to \mathcal{B}
2. Sample a mini-batch $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j\}$ from \mathcal{B} uniformly
3. Compute $y_j = r(\mathbf{s}_j, \mathbf{a}_j) + \gamma \text{softmax}_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mathbf{a}'_j)$ using the target network $Q_{\phi'}$
4. $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_{\phi}}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_{\phi}(\mathbf{s}_j, \mathbf{a}_j) - y_j)$
5. Update ϕ' : copy ϕ every N steps or using Polyak averaging $\phi' \leftarrow \tau\phi' - (1 - \tau)\phi$

end

14 Inverse Reinforcement Learning

14.1 Algorithms

14.1.1 Maximum Entropy Inverse Reinforcement Learning

Algorithm 25: MaxEnt IRL

Loop

1. Given ψ , compute backward message $\beta(\mathbf{s}_t, \mathbf{a}_t) = p(\mathcal{O}_{t:T} | \mathbf{s}_t, \mathbf{a}_t)$
2. Given ψ , compute forward message $\alpha(\mathbf{s}_t) = p(\mathbf{s}_t | \mathcal{O}_{1:t-1})$
3. Compute $\mu_t(\mathbf{s}_t, \mathbf{a}_t) \propto \beta(\mathbf{s}_t, \mathbf{a}_t) \alpha(\mathbf{s}_t)$
4. Evaluate
$$\nabla_{\psi} \mathcal{L} = \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\psi} r_{\psi}(\mathbf{s}_t^i, \mathbf{a}_t^i) - \sum_{t=1}^T \iint \mu_t(\mathbf{s}_t, \mathbf{a}_t) \nabla_{\psi} r_{\psi}(\mathbf{s}_t, \mathbf{a}_t) d\mathbf{s}_t d\mathbf{a}_t$$
5. $\psi \leftarrow \psi + \eta \nabla_{\psi} \mathcal{L}$

end

14.2 Tips & hacks

1. Loss computation:

- $\nabla_{\psi} \mathcal{L} \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\psi} r_{\psi}(\tau_i) - \frac{1}{M} \sum_{j=1}^M \nabla_{\psi} r_{\psi}(\tau_j)$
- With importance sampling: $\nabla_{\psi} \mathcal{L} \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\psi} r_{\psi}(\tau_i) - \frac{1}{\sum_j w_j} \sum_{j=1}^M w_j \nabla_{\psi} r_{\psi}(\tau_j)$
where $w_j = \frac{p(\tau) \exp r_{\psi}(\tau_j)}{\pi(\tau_j)} = \frac{\exp \sum_t r_{\psi}(\mathbf{s}_t, \mathbf{a}_t)}{\prod_t \pi(\mathbf{a}_t | \mathbf{s}_t)}$

2. IRL can be treated as a GAN

15 Transfer & Multi-task Learning

15.1 Tips & hacks

1. Forward transfer: train on one task, transfer to a new task

- Finetune
- Randomize the source domains

2. Multi-task transfer: train on many tasks, transfer to a new task

- Generate highly randomized source domains
- Model-based RL
- Model distillation
- Contextual policies
- Modular policy networks

3. Multi-task meta-learning: learn to learn from many tasks

- RNN-based meta-learning
- Gradient-based meta-learning

16 Distributed Reinforcement Learning

17 Exploration

17.1 Algorithms

17.1.1 Pre-train & finetune

Algorithm 26: Pre-train & finetune

1. Collect demonstration data $(\mathbf{s}_i, \mathbf{a}_i)$
 2. Initialize π_θ as $\max_\theta \sum_i \log \pi_\theta(\mathbf{a}_i | \mathbf{s}_i)$
- Loop**
3. Run π_θ to collect experience
 4. Improve π_θ with any RL algorithm
- end**
-

17.2 Tips & hacks

1. Exploration bonus: $r^+(\mathbf{s}, \mathbf{a}) + \mathcal{B}(N(\mathbf{s}))$

- UCB: $\mathcal{B}(N(\mathbf{s})) = \sqrt{\frac{2 \ln n}{N(\mathbf{s})}}$
- MBIE-EB: $\mathcal{B}(N(\mathbf{s})) = \frac{1}{\sqrt{N(\mathbf{s})}}$
- BEB: $\mathcal{B}(N(\mathbf{s})) = \frac{1}{N(\mathbf{s})}$
- Use pseudo-counts and GANs
- Counts encoding into hashes
- Density modeling
- Heuristic estimation of counts

2. Thomson sampling

3. Bootstrap

4. Information gain

5. Data initialization: initialize data (e.g for Q-learning) with demonstrations as off-policy data

- Use good and bad demonstrations

- Use importance sampling

6. **Hybrid objective:** $\underbrace{E_{\pi_{\theta}}[r(\mathbf{s}, \mathbf{a})]}_{RL} + \lambda \underbrace{\sum_{(\mathbf{s}, \mathbf{a}) \in \mathcal{D}_{demos}} \log \pi_{\theta}(\mathbf{a}|\mathbf{s})}_{IL}$

18 Meta-learning

19 Information theory